



Week 3 Assignment: Implement a Quadratic Layer

In this week's programming exercise, you will build a custom quadratic layer which computes $y = ax^2 + bx + c$. Similar to the ungraded lab, this layer will be plugged into a model that will be trained on the MNIST dataset. Let's get started!

Imports

```
In [1]: import tensorflow as tf
        from tensorflow.keras.layers import Layer

        import utils
```

Define the quadratic layer (TODO)

Implement a simple quadratic layer. It has 3 state variables: a , b and c . The computation returned is $ax^2 + bx + c$. Make sure it can also accept an activation function.

`__init__`

- call `super(my_fun, self)` to access the base class of `my_fun`, and call the `__init__()` function to initialize that base class. In this case, `my_fun` is `SimpleQuadratic` and its base class is `Layer`.
- `self.units`: set this using one of the function parameters.
- `self.activation`: The function parameter `activation` will be passed in as a string. To get the tensorflow object associated with the string, please use `tf.keras.activations.get()`

`build`

The following are suggested steps for writing your code. If you prefer to use fewer lines to implement it, feel free to do so. Either way, you'll want to set `self.a`, `self.b` and `self.c`.

- `a_init`: set this to `tensorflow's random_normal_initializer()`
- `a_init_val`: Use the `random_normal_initializer()` that you just created and invoke it, setting the `shape` and `dtype`.
 - The `shape` of `a` should have its row dimension equal to the last dimension of `input_shape`, and its column dimension equal to the number of units in the layer.
 - This is because you'll be matrix multiplying $x^2 * a$, so the dimensions should be compatible.
 - set the `dtype` to `'float32'`
- `self.a`: create a tensor using `tf.Variable`, setting the `initial_value` and set trainable to `True`.
- `b_init`, `b_init_val`, and `self.b`: these will be set in the same way that you implemented `a_init`, `a_init_val` and `self.a`
- `c_init`: set this to `tf.zeros_initializer()`
- `c_init_val`: Set this by calling the `tf.zeros_initializer` that you just instantiated, and set the `shape` and `dtype`
 - `shape`: This will be a vector equal to the number of units. This expects a tuple, and remember that a tuple `(9,)` includes a comma.
 - `dtype`: set to `'float32'`.
- `self.c`: create a tensor using `tf.Variable`, and set the parameters `initial_value` and `trainable`.

`call`

The following section performs the multiplication $x^2a + xb + c$. The steps are broken down for clarity, but you can also perform this calculation in fewer lines if you prefer.

- `x_squared`: use `tf.math.square()`
- `x_squared_times_a`: use `tf.matmul()`.
 - If you see an error saying `InvalidArgumentError: Matrix size-incompatible`, please check the order of the matrix multiplication to make sure that the matrix dimensions line up.
- `x_times_b`: use `tf.matmul()`.
- `x2a_plus_xb_plus_c`: add the three terms together.
- `activated_x2a_plus_xb_plus_c`: apply the class's `activation` to the sum of the three terms.

```
In [2]: ## Please uncomment all lines in this cell and replace those marked with `# YOUR CODE HERE`.
        ## You can select all lines in this code cell with Ctrl+A (Windows/Linux) or Cmd+A (Mac), then press Ctrl+/ (Windows/Linux) c

class SimpleQuadratic(Layer):

    def __init__(self, units=32, activation=None):
        '''Initializes the class and sets up the internal variables'''
        # YOUR CODE HERE
        super(SimpleQuadratic, self).__init__()
        self.units = units
        self.activation = tf.keras.activations.get(activation)

    def build(self, input_shape):
        '''Create the state of the layer (weights)'''
        # a and b should be initialized with random normal, c (or the bias) with zeros.
        # remember to set these as trainable.
        # YOUR CODE HERE
        a_init = tf.random_normal_initializer()
        b_init = tf.random_normal_initializer()
        c_init = tf.zeros_initializer()

        self.a = tf.Variable(name="kernel", initial_value = a_init(shape= (input_shape[-1], self.units),
                                                                    dtype= "float32"), trainable = True)

        self.b = tf.Variable(name="kernel", initial_value = b_init(shape= (input_shape[-1], self.units),
                                                                    dtype= "float32"), trainable = True)

        self.c = tf.Variable(name="bias", initial_value = c_init(shape= (self.units,),
                                                                    dtype= "float32"), trainable = True)

    def call(self, inputs):
        '''Defines the computation from inputs to outputs'''
        # YOUR CODE HERE
        result = tf.matmul(tf.math.square(inputs), self.a) + tf.matmul(inputs, self.b) + self.c
        return self.activation(result)
```

Test your implementation

```
In [3]: ▶ utils.test_simple_quadratic(SimpleQuadratic)
```

All public tests passed

Train your model with the `SimpleQuadratic` layer that you just implemented.

```
In [4]: ▶ # THIS CODE SHOULD RUN WITHOUT MODIFICATION
# AND SHOULD RETURN TRAINING/TESTING ACCURACY at 97%+
```

```
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    SimpleQuadratic(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11493376/11490434 [=====] - 0s 0us/step
Train on 60000 samples
Epoch 1/5
60000/60000 [=====] - 12s 204us/sample - loss: 0.2687 - accuracy: 0.9208
Epoch 2/5
60000/60000 [=====] - 12s 197us/sample - loss: 0.1306 - accuracy: 0.9616
Epoch 3/5
60000/60000 [=====] - 12s 198us/sample - loss: 0.0980 - accuracy: 0.9699
Epoch 4/5
60000/60000 [=====] - 12s 198us/sample - loss: 0.0836 - accuracy: 0.9738
Epoch 5/5
60000/60000 [=====] - 12s 198us/sample - loss: 0.0711 - accuracy: 0.9772
10000/10000 [=====] - 1s 77us/sample - loss: 0.0745 - accuracy: 0.9778
```

```
Out[4]: [0.07447561496943235, 0.9778]
```

```
In [ ]: ▶
```